

Figure 8 - Speculative decoding speedup (Leviathan, ICML 2023)

